


























Predicting Flight Delays Using SparkML on AWS

Jason Carpenter | Christopher Csiszar | Nishan Madawanarachchi | Ryan Campa

After a long weekend in Vegas...



 DEPARTURES						
FLIGHT	DESTINATION		TIME	STATUS	GATE	
PA 0030	SAN FRANCISCO		00:30	BOARDING	12	 
LX 3456	LONDON		01:45	GO TO GATE	34	 
BA 0300	SINGAPORE		02:15	ON TIME	15	 
LA 0200	LOS ANGELES		02:00	CANCELLED	13	 
FE 4561	BRISBANE		02:30	ON TIME	17	 
LX 4523	LONDON HTHROW		02:45	ON TIME	19	 
BK 4313	MEXICO CITY		02:30	ON TIME	25	 
HI 6123	UPSPOUP1111		02:30	ON TIME	09	 
QW 1173	MANILA		03:00	ON TIME	26	 
CQ 2123	CHICAGO		03:25	ON TIME	21	 
BX 0067	PORTLAND		03:30	ON TIME	04	 
WA 0264	BALI		03:45	ON TIME	03	 

Data



- DOT's Bureau of Transportation and Statistics
- Tracks on-time performance of domestic flights operated by large air carriers.
- Early, On time, delayed, diverted, cancelled
- 2006-2008
- CSV, ~2Gb, 21 million rows, 29 columns
- <http://stat-computing.org/dataexpo/2009/the-data.html>

```
In [3]: data_2007.transpose().iloc[:,0]
```

```
Out[3]: Year                2007  
Month                    1  
DayOfMonth               1  
DayOfWeek                1  
DepTime                 1232  
CRSDepTime              1225  
ArrTime                 1341  
CRSArrTime              1340  
UniqueCarrier           WN  
FlightNum               2891  
TailNum                 N351  
ActualElapsedTime       69  
CRSElapsedTime          75  
AirTime                 54  
ArrDelay                1  
DepDelay                7  
Origin                  SMF  
Dest                    ONT  
Distance                389  
TaxiIn                  4  
TaxiOut                 11  
Cancelled                0  
CancellationCode       NaN  
Diverted                0  
CarrierDelay            0  
WeatherDelay            0  
NASDelay                0  
SecurityDelay           0  
LateAircraftDelay       0  
Name: 0, dtype: object
```

Objective



- Build Machine Learning Pipeline on AWS using Apache Spark and MongoDB
- Predict Arrival Delays
- Evaluate Performance of Model

Pipeline



Data Storage



aws

Services ▾ Resource Groups ▾

Jason Carpenter ▾ Global ▾ Support ▾

Amazon S3 > msan697-flightdelays

Overview

Properties

Permissions Public

Management

Q Type a prefix and press Enter to search. Press ESC to clear.

Upload

Create folder

More ▾

US West (Oregon) ↻

Viewing 1 to 3				
<input type="checkbox"/>	Name ↑ ▾	Last modified ↑ ▾	Size ↑ ▾	Storage class ↑ ▾
<input type="checkbox"/>	2006.csv	Jan 10, 2018 4:32:58 PM GMT-0800	640.9 MB	Standard
<input type="checkbox"/>	2007.csv	Jan 10, 2018 4:43:34 PM GMT-0800	670.3 MB	Standard
<input type="checkbox"/>	2008.csv	Jan 10, 2018 4:25:41 PM GMT-0800	657.5 MB	Standard

Viewing 1 to 3

AWS Configuration



- **m3.2xlarge** spark-ec2 stand-alone cluster with 1 master and 2 workers

VCPU	ECU	Memory(GiB)	Instance Storage (GB)
8	26	30	2 x 80 (SSD)

MongoDB



```
> use msan697
switched to db msan697
> db.flightdelays.findOne({'Year': 2007})
{
  "_id" : ObjectId("5a58289ac93ee71563909487"),
  "Year" : 2007,
  "Month" : 1,
  "DayOfMonth" : 1,
  "DayOfWeek" : 1,
  "DepTime" : 1232,
  "CRSDepTime" : 1225,
  "ArrTime" : 1341,
  "CRSArrTime" : 1340,
  "UniqueCarrier" : "WN",
  "FlightNum" : 2891,
  "TailNum" : "N351",
  "ActualElapsedTime" : 69,
  "CRSElapsedTime" : 75,
  "AirTime" : 54,
  "ArrDelay" : 1,
  "DepDelay" : 7,
  "Origin" : "SMF",
  "Dest" : "ONT",
  "Distance" : 389,
  "TaxiIn" : 4,
  "TaxiOut" : 11,
  "Cancelled" : 0,
  "CancellationCode" : "",
  "Diverted" : 0,
  "CarrierDelay" : 0,
  "WeatherDelay" : 0,
  "NASDelay" : 0,
  "SecurityDelay" : 0,
  "LateAircraftDelay" : 0
}
```


SparkSQL



```
>>> df = spark.read.format("com.mongodb.spark.sql.DefaultSource").option('uri','mongodb://127.0.0.1/msan697.flightdelays').load()
18/01/18 21:06:02 WARN MongoInferSchema: Field 'ActualElapsedTime' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'AirTime' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'ArrDelay' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'ArrTime' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'CarrierDelay' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'DepDelay' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'DepTime' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'LateAircraftDelay' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'NASDelay' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'SecurityDelay' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'TailNum' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'TaxiIn' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'TaxiOut' contains conflicting types converting to StringType
18/01/18 21:06:02 WARN MongoInferSchema: Field 'WeatherDelay' contains conflicting types converting to StringType
>>> df.printSchema()
root
|-- ActualElapsedTime: string (nullable = true)
|-- AirTime: string (nullable = true)
|-- ArrDelay: string (nullable = true)
|-- ArrTime: string (nullable = true)
|-- CRSArrTime: integer (nullable = true)
|-- CRSDepTime: integer (nullable = true)
|-- CRSElapsedTime: integer (nullable = true)
|-- CancellationCode: string (nullable = true)
|-- Cancelled: integer (nullable = true)
|-- CarrierDelay: string (nullable = true)
|-- DayOfWeek: integer (nullable = true)
|-- DayOfMonth: integer (nullable = true)
|-- DepDelay: string (nullable = true)
|-- DepTime: string (nullable = true)
|-- Dest: string (nullable = true)
|-- Distance: integer (nullable = true)
|-- Diverted: integer (nullable = true)
|-- FlightNum: integer (nullable = true)
|-- LateAircraftDelay: string (nullable = true)
|-- Month: integer (nullable = true)
|-- NASDelay: string (nullable = true)
|-- Origin: string (nullable = true)
|-- SecurityDelay: string (nullable = true)
|-- TailNum: string (nullable = true)
|-- TaxiIn: string (nullable = true)
|-- TaxiOut: string (nullable = true)
|-- UniqueCarrier: string (nullable = true)
|-- WeatherDelay: string (nullable = true)
|-- Year: integer (nullable = true)
|-- _id: struct (nullable = true)
|   |-- oid: string (nullable = true)
```

Creating Data Frame From MongoDB

SparkSQL



```
>>> df_min = df.filter('Cancelled = 0 and Diverted = 0').select('Year','Month',  
'DayOfMonth','DayOfWeek','CRSDepTime','CRSArrTime','Distance','ArrDelay','CRSEla  
psedTime','UniqueCarrier','FlightNum','TailNum','Origin','Dest')  
>>> df_min.printSchema()  
root  
|-- Year: integer (nullable = true)  
|-- Month: integer (nullable = true)  
|-- DayOfMonth: integer (nullable = true)  
|-- DayOfWeek: integer (nullable = true)  
|-- CRSDepTime: integer (nullable = true)  
|-- CRSArrTime: integer (nullable = true)  
|-- Distance: integer (nullable = true)  
|-- ArrDelay: string (nullable = true)  
|-- CRSElapsedTime: integer (nullable = true)  
|-- UniqueCarrier: string (nullable = true)  
|-- FlightNum: integer (nullable = true)  
|-- TailNum: string (nullable = true)  
|-- Origin: string (nullable = true)  
|-- Dest: string (nullable = true)
```

Filtered out cancelled and diverted flights and removed data leakage columns

SparkSQL



Using SparkSQL, we converted strings to category codes

```
>>> sqlContext.sql('select * from flightsnumeric').show()
18/01/18 21:12:05 WARN SizeEstimator: Failed to check whether UseCompressedOops is set; assuming yes
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Year|Month|DayofMonth|DayOfWeek|CRSDepTime|CRSArrTime|Distance|ArrDelay|CRSElapsedTime|UniqueCarrier|FlightNum|TailNum|Origin|Dest|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|2008|1|16|3|1910|2130|375|-18|80|16.0|3428.0|4461.0|12.0|164.0|
|2008|1|5|6|1105|1340|490|-9|95|16.0|2917.0|4630.0|144.0|12.0|
|2008|1|18|5|1910|2130|375|244|80|16.0|3428.0|4630.0|12.0|164.0|
|2008|1|19|6|1910|2130|375|-1|80|16.0|3428.0|4495.0|12.0|164.0|
|2008|1|20|7|1910|2130|375|152|80|16.0|3428.0|4702.0|12.0|164.0|
|2008|1|21|1|1910|2130|375|28|80|16.0|3428.0|5017.0|12.0|164.0|
|2008|1|24|4|1910|2130|375|35|80|16.0|3428.0|4666.0|12.0|164.0|
|2008|1|25|5|1910|2130|375|0|80|16.0|3428.0|4527.0|12.0|164.0|
|2008|1|22|2|1910|2130|375|2|80|16.0|3428.0|4633.0|12.0|164.0|
|2008|1|26|6|1910|2130|375|-6|80|16.0|3428.0|4653.0|12.0|164.0|
|2008|1|27|7|1910|2130|375|-8|80|16.0|3428.0|4679.0|12.0|164.0|
|2008|1|28|1|1910|2130|375|27|80|16.0|3428.0|4716.0|12.0|164.0|
|2008|1|30|3|1910|2130|375|14|80|16.0|3428.0|4667.0|12.0|164.0|
|2008|1|7|1|1600|1746|449|28|106|16.0|3428.0|4489.0|31.0|12.0|
|2008|1|8|2|1600|1746|449|15|106|16.0|3428.0|5595.0|31.0|12.0|
|2008|1|9|3|1600|1746|449|8|106|16.0|3428.0|4549.0|31.0|12.0|
|2008|1|10|4|1600|1746|449|17|106|16.0|3428.0|4434.0|31.0|12.0|
|2008|1|29|2|1910|2130|375|32|80|16.0|3428.0|4489.0|12.0|164.0|
|2008|1|12|6|1600|1746|449|1|106|16.0|3428.0|4660.0|31.0|12.0|
|2008|1|13|7|1600|1746|449|-16|106|16.0|3428.0|4640.0|31.0|12.0|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

Querying Data Using SparkSQL




- Feature Engineering:
 - Odd/Even Flight Num
 - odd flight numbers correspond to westbound/southbound flights
 - even flight numbers correspond to eastbound/northbound flights
 - Num Digits of Flight Num
 - Less than 3 digits correspond to long-haul or otherwise premium flights

SparkML



```
best_rf = RandomForestRegressor(numTrees=40,  
                                maxDepth=15,  
                                minInstancesPerNode=80,  
                                maxBins=512,  
                                subsamplingRate=0.5)  
best_rfmodel = best_rf.fit(trn)  
  
preds = best_rfmodel.transform(test)  
eval = RegressionEvaluator(metricName='rmse')  
rmse = eval.evaluate(preds)  
  
print('RMSE = %.4f' % rmse)  
preds.toPandas().to_csv('rf_regression_preds.csv')
```

Results



Oct-Dec 2008 Flight Delay Prediction Errors - RMSE = 12.4807



Lessons Learned



- Having undesirable number of partitions led the machine learning code to run 3x slower
- Running Random Forest Regression with 3 m1.large instances (default machines in ec2-cluster) led to memory issues.
- Diagnose code issues by first running only a fraction of your data, not all of it.
- Due to time constraints, we used maxDepth=15, preventing our model from capturing the high cardinality of the Origin/Destination features